
PaCX-MAE: Physiology-Augmented Chest X-Ray Masked Autoencoder

Yancheng Liu¹ Kenichi Maeda¹ Manan Pancholy¹

Abstract

Clinical diagnosis often requires combining imaging with physiological measurements, yet deployed models typically operate on unimodal data. We present **PaCX-MAE**, a cross-modal distillation framework that injects physiological priors into chest X-ray (CXR) encoders while remaining strictly unimodal at inference. PaCX-MAE augments in-domain masked autoencoding with a dual contrastive-predictive objective, aligning CXR representations with paired ECG and laboratory embeddings. Extensive evaluation across nine benchmarks demonstrates consistent improvements over domain-specific MAE, particularly on physiology-dependent tasks (e.g., **+2.7 AUROC** on MedMod; **+6.5 F1** on VinDr). The method proves highly label-efficient in the **1%** regime and preserves anatomical fidelity, achieving parity with MAE on segmentation tasks. Zero-shot and attention analyses confirm that PaCX-MAE successfully learns to attend to physiological indicators, such as the cardiac silhouette, absent in standard visual pretraining.

1. Introduction

Self-supervised learning on medical imaging typically considers modalities in isolation or assumes full modality availability at inference. While multimodal models can leverage physiological signals (e.g., ECG, PPG, labs) to capture dynamic patient states such as cardiac activity or fluid balance, this assumption frequently breaks in clinical practice. In acute settings, a Chest X-ray is often the only immediately available modality. Consequently, models trained on rich multimodal data fail to deploy because the auxiliary modalities are missing, while unimodal models fail to capture the systemic physiological context invisible to the naked eye.

¹Department of Computer Science, Brown University, Providence, RI, USA. Correspondence to: Yancheng Liu <yancheng_liu@brown.edu>.

Assignment created by Randall Balastriero for CSCI2952X: Research Topics in Self-Supervised Learning, offered in Fall 2025. Copyright 2025 by the author.

However, physiological states manifest as latent “fingerprints” in anatomical imaging—for instance, vascular redistribution signaling fluid overload. This raises a critical question: *Can a vision encoder learn to infer physiological context from anatomy alone?* We hypothesize that aligning visual features with physiological encodings during training enables a model to recognize these subtle anatomical correlates without requiring auxiliary inputs at inference.

To this end, we introduce **PaCX-MAE**, a framework that distills physiological priors into a strictly unimodal vision encoder. Unlike standard multimodal fusion, PaCX-MAE uses paired data only during pretraining to align CXR representations with physiological embeddings via a dual contrastive-predictive objective. We validate PaCX-MAE across nine benchmarks, showing that it: (1) outperforms standard unimodal MAE, particularly on physiology-heavy tasks (e.g., MedMod); (2) maintains pixel-level fidelity for segmentation; and (3) significantly improves label efficiency in low-data regimes.

2. Related Work

SSL for Medical Imaging Self-supervised learning in medical imaging has largely relied on contrastive methods, which learn invariance by maximizing agreement between augmented views (Azizi et al., 2021; Cho et al., 2023; Gorade et al., 2025). However, recent work indicates that Masked Autoencoders (MAE) often yield superior representations for downstream diagnosis by forcing the reconstruction of fine-grained anatomical details rather than global invariants (Xiao et al., 2023; Zhou et al., 2023; Huang et al., 2023). **PaCX-MAE** takes advantage of this reconstructive strength, initializing with an MAE backbone to ensure robust anatomical features before injecting physiological priors.

Multimodal Distillation & Missing Modalities While standard multimodal learning fuses data streams (e.g., imaging, text, continuous waveforms) during both training and inference (Zhang et al., 2022; Radford et al., 2021), clinical deployment is often constrained by missing modalities. To address this, cross-modal knowledge distillation transfers “privileged information” from a multimodal teacher to a unimodal student (Lopez-Paz et al., 2016; Gupta et al., 2015). In medical domains, this paradigm has been applied to distill radiology reports (Tiu et al., 2022; Boecking et al., 2022) or

missing MRI sequences (Dou et al., 2020; Wang et al., 2023) into image-only encoders. **PaCX-MAE** extends this framework to physiology, distilling dense signals (ECG, labs) into CXR representations to enable physiological reasoning at inference time without requiring auxiliary sensors.

3. Methodology

We present **PaCX-MAE**, a framework for distilling physiological priors into a unimodal vision backbone. As illustrated in Figure 1, our approach decouples *representation learning* from *cross-modal alignment* via a two-stage curriculum: (1) independent unimodal pretraining to establish robust feature spaces, followed by (2) cross-modal distillation that aligns the visual encoder with frozen physiological targets via a dual contrastive-predictive objective.

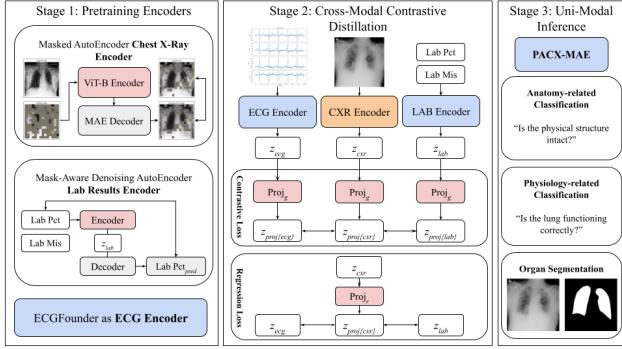


Figure 1. Overview of the PaCX architecture. The pipeline comprises unimodal pretraining (Stage 1) and cross-modal distillation (Stage 2). Colors indicate optimization status: red (trainable), orange (LoRA-adapted), and blue (frozen). During distillation, the CXR encoder learns to predict physiological embeddings via lightweight heads, which are discarded at inference.

3.1. Stage 1: Unimodal Pretraining

We utilize the Symile-MIMIC dataset (Saporta et al., 2025), containing $N \approx 10k$ paired triplets of CXR (x_v), ECG (x_e), and Laboratory (x_l) data. To mitigate the risk of overfitting to limited paired data, we first initialize modality-specific encoders on large-scale external datasets.

Vision Encoder (CXR). We initialize the vision backbone (f_v) using a ViT-B architecture trained via Masked Autoencoding (MAE) on CheXpert (Irvin et al., 2019). We employ an aggressive masking ratio of 0.90, as it forces the model to infer global anatomical semantics (e.g., cardiac silhouette, mediastinal width) from limited visual cues rather than exploiting local texture shortcuts (Gupta et al., 2024). We prioritize this reconstructive objective over contrastive alternatives (e.g., MoCo, DINO). Unlike contrastive methods, which enforce invariance to augmentations that can obscure subtle intensity cues (e.g., fluid opacity) (Huang et al., 2023), MAE preserves fine-grained details critical for physiological inference and avoids the semantic collapse often associated

with false-negative sampling in medical imaging (see Fig. 2). Implementation details, including the specific optimizer and scheduler configurations, are provided in Appendix A.1.1.



Figure 2. MAE Pretraining Reconstructions. Top: Original CXRs; Bottom: Reconstructions under 90% masking. Despite extreme sparsity, the model accurately recovers key physiological indicators such as the cardiac boundary and diaphragm curvature.

Physiological Targets (ECG & Labs). We employ high-fidelity, frozen encoders to serve as distillation targets. For Laboratory data, we pretrain a **mask-aware Denoising Autoencoder** that models both measured values and structured missingness typical of tabular clinical records. For ECG, we utilize **ECGFounder** (Li et al., 2025), a transformer pre-trained on 10 million recordings to capture high-frequency morphological patterns. Crucially, these encoders (f_e, f_l) remain **frozen** during Stage 2. This prevents degenerate co-adaptation and ensures that physiological structure is transferred into the visual manifold rather than jointly relearned. Architectural and training details are provided in Appendix A.1.1.

3.2. Stage 2: Cross-Modal Distillation

The objective of Stage 2 is to inject physiological priors into the visual backbone f_v without degrading its anatomical fidelity. We employ a dual-branch architecture that aligns the CXR embeddings with the frozen physiological targets (f_e, f_l) via a hybrid contrastive-regression loss.

Parameter-Efficient Adaptation (LoRA). To prevent catastrophic forgetting of the dense anatomical priors learned during MAE pretraining, we freeze the majority of the visual backbone, keeping only the normalization layers trainable to stabilize feature distributions. We inject Low-Rank Adaptation (LoRA) matrices into both the attention ('qkv') and feed-forward ('fc') modules. This ensures that the core visual manifold remains stable, while the lightweight LoRA parameters (< 1% of total weights) learn the specific projections required to bridge the modality gap.

Dual Distillation Objective. We define a hybrid objective $\mathcal{L}_{total} = \lambda_C \mathcal{L}_{contrastive} + \lambda_R \mathcal{L}_{regression}$ to balance global semantic alignment with latent feature reconstruction.

1. **Global Contrastive Alignment (\mathcal{L}_C):** We project CXR, ECG, and Lab embeddings into a shared latent space and align them via a symmetric InfoNCE loss with a learnable temperature parameter. To prevent the model from exploiting local batch biases, we employ **global negative sampling**, gathering embeddings across all dis-

tributed GPUs to maximize the diversity of negative pairs. We further apply label smoothing ($\epsilon = 0.02$) to prevent overfitting to noisy medical labels.

- 2. Latent Regression (\mathcal{L}_R):** Contrastive learning alone can settle for “shortcut” features sufficient for discrimination but insufficient for detailed reasoning. To counter this, we introduce modality-specific regression heads that predict the *exact* unprojected frozen embedding vectors of the physiological encoders. We minimize the **Cosine Distance** ($1 - \cos(\hat{y}, y)$) between the predicted and target embeddings, forcing the visual encoder to internalize the dense semantic structure of the physiological signals.

Both the projection (contrastive) and regression heads are discarded at inference, leaving only the enhanced visual encoder. Implementation details are provided in Appendix A.1.2.

4. Experiments

Datasets and Baselines. We evaluate PaCX-MAE (also denoted as PaCX in short) across a comprehensive suite of 9 public benchmarks covering binary, multiclass, and multi-label classification, as well as semantic segmentation tasks. Detailed descriptions and statistics for each dataset are provided in Appendix A.2.1. We compare our method against two primary baselines: a standard **ImageNet-pretrained** ViT-B/16 and a domain-specific **Masked Autoencoder (MAE)** pretrained on unimodal CXRs. Since PaCX utilizes the same backbone architecture as the MAE baseline, any performance variance directly isolates the efficacy of our cross-modal physiological distillation.

Evaluation Protocols. Our evaluation strategy assesses representation quality through **linear probing**, where a linear classifier is trained on top of the frozen encoder backbone. We rigorously test **data efficiency** by training on restricted subsets (e.g., 1%, 10%) of the available training data. To deconstruct the impact of specific physiological signals, we conduct **modality ablations** and loss component analyses. Furthermore, we evaluate the alignment of learned representations via **zero-shot retrieval** metrics (Recall@K, Cosine Similarity) and provide qualitative interpretations using **Attention Rollout**. Complete implementation details, including hyperparameters, optimization schedules, and ablation protocols, are provided in Appendix A.2.2.

5. Results

5.1. Clinical Transfer & Data Efficiency

Domain-Specific Pretraining. Table 1 confirms that domain-specific pretraining consistently outperforms ImageNet initialization. Crucially, PaCX retains the dense anatomical competence of the strong MAE baseline, achieving parity on pixel-precise segmentation benchmarks like

CXL-Seg (0.996 IoU) and **COVID-QU-Ex (0.942 IoU)**, demonstrating that physiological distillation does not induce catastrophic forgetting of structural features. We observe a minor drop on QaTa-COV19 (−1.1 IoU), likely due to domain-specific infection patterns not captured by our physiological priors; this does not affect overall conclusions.

Physiological Awareness. PaCX demonstrates superior transfer to tasks reflecting latent physiological states rather than purely visible structures. On physiology-dense benchmarks—**CheXchoNet**, **VinDr-CXR**, and **MedMod**—it significantly outperforms the unimodal MAE, achieving gains of **+2.4 AUROC** and **+6.5 F1** on VinDr-CXR, **+2.7 AUROC** on MedMod, and **+5.1 F1** on CheXchoNet. This supports our hypothesis that distilled “phantom” physiological signals enhance diagnosis where visual cues are subtle.

Dataset	Metric	ImageNet	MAE	PaCX	Δ (pp)
TB ¹	AUROC	0.887	0.899	0.910	+1.1
	F1	0.818	0.814	0.846	+3.2
CheXchoNet ¹	AUROC	0.728	0.788	0.803	+1.5
	F1	0.147	0.215	0.266	+5.1
ChestX6 ²	AUROC	0.983	0.988	0.989	+0.1
	F1	0.876	0.905	0.906	+0.1
VinDr-CXR ³	AUROC	0.751	0.847	0.871	+2.4
	F1	0.097	0.191	0.256	+6.5
NIH-14 ³	AUROC	0.721	0.772	0.783	+1.1
	F1	0.048	0.113	0.115	+0.2
MedMod ³	AUROC	0.612	0.695	0.722	+2.7
	F1	0.091	0.231	0.253	+2.2
COVID-QU-Ex ⁴	IoU	0.894	0.942	0.942	0.0
	Dice	0.943	0.970	0.970	0.0
QaTa-COV19 ⁴	IoU	0.622	0.726	0.715	-1.1
	Dice	0.766	0.841	0.833	-0.8
CXL-Seg ⁴	IoU	0.984	0.996	0.996	0.0
	Dice	0.992	0.998	0.998	0.0

Table 1. Comparison of classification and segmentation performance with different pretraining methods. Superscripts denote task type: ¹ binary, ² multiclass, ³ multilabel, ⁴ segmentation.

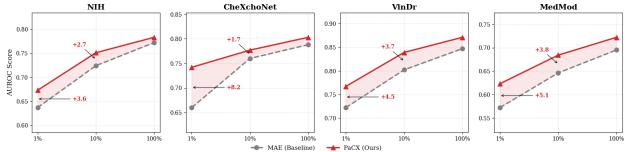


Figure 3. Label Efficiency. PaCX (red) outperforms MAE (grey) consistently at 1% and 10% training data, demonstrating robust few-shot transfer.

Low-Data Efficiency. PaCX significantly lowers sample complexity. As illustrated in Figure 3, the performance gap is widest in extremely low-data regimes. In the **1% data regime**, PaCX consistently surpasses the MAE baseline, showing AUROC improvements of **+8.2** on CheXchoNet, and $\sim+5$ on MedMod and VinDr. Even at **10% data**, it maintains a robust lead of **+1.7–+3.8** across benchmarks. This consistent advantage indicates that learned physiological priors act as effective regularizers.

5.2. Physiological Alignment & Interpretability

Zero-Shot Alignment. Table 2 confirms that PaCX effectively distills physiological signals, achieving superior structural alignment between frozen CXR embeddings and ground-truth targets. PaCX surpasses the MAE baseline in Cosine Similarity across both modalities (**0.229 ECG, 0.252 Labs**), indicating it successfully pulls anatomically distinct X-rays closer to their correct physiological profiles in the latent space without explicit supervision.

Metric	ECG Targets			Lab Targets		
	ImNet	MAE	PaCX	ImNet	MAE	PaCX
Cos Sim	0.143	0.204	0.229	0.187	0.239	0.252
R@5	1.51%	5.17%	5.60%	1.51%	3.66%	3.66%

Table 2. Zero-Shot Alignment. PaCX improves latent structure across both modalities, demonstrating superior or equivalent retrieval capability (R@5) and vector alignment (Cosine Similarity).

Attention Rollout. We visualize the impact of this alignment using Attention Rollout on validation “rescue cases”—instances where PaCX correctly classified pathology that the baseline missed. Figure 4 illustrates a representative case of cardiomegaly. While the MAE baseline scatters attention over bony structures like clavicles, PaCX tightens focus on the cardiac silhouette and mediastinum, as highlighted by the red regions in the Difference Map. This confirms that physiological signals effectively guide the visual encoder toward clinically relevant soft-tissue anatomy. Additional cases in Appendix A.3.1 confirm this is a systematic trend, with PaCX consistently exhibiting more focused attention on relevant organ systems.

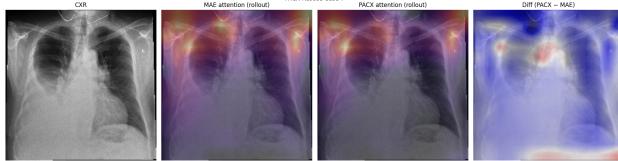


Figure 4. Attention Shift. PaCX (middle-right) shifts focus from bony structures to the cardiac silhouette (red in Difference Map).

5.3. Component Analysis

To disentangle the contributions of specific signals and objectives, we analyze component-wise performance on three physiology-dense benchmarks (Table 3).

Modality Synergy. Each modality helps on at least one dataset/metric; the full PaCX configuration provides the most balanced trade-off across tasks. It confirms that ECG and Laboratory data provide complementary rather than redundant signals. For instance, removing Lab priors causes a notable drop in F1 on MedMod ($0.253 \rightarrow 0.243$) and VinDr ($0.256 \rightarrow 0.233$).

Objective Function. Comparing loss components reveals that regression alone (\mathcal{L}_{reg}) is insufficient for robust representation learning, significantly underperforming the contrastive objective (\mathcal{L}_{cont}). However, the hybrid objective ($\mathcal{L}_{cont} + \mathcal{L}_{reg}$) achieves the best overall stability (e.g., VinDr F1 +1.5 over \mathcal{L}_{cont} alone). This loss ablation study validates the design choice to combine global structural alignment with dense continuous regression.

Dataset	Metric	Modality Ablation			Loss Ablation		
		ECG	Lab	PaCX	Cont	Reg	PaCX
CheXchoNet	AUC	0.801	0.795	0.803	0.799	0.789	0.803
	F1	0.296	0.275	0.266	0.273	0.227	0.266
MedMod	AUC	0.717	0.721	0.722	0.722	0.673	0.722
	F1	0.243	0.245	0.253	0.258	0.131	0.253
VinDr	AUC	0.871	0.875	0.871	0.866	0.843	0.871
	F1	0.233	0.248	0.256	0.241	0.130	0.256

Table 3. Ablation Study. Comparing single-modality sources (Left) and isolated loss components (Right). The full PaCX configuration yields the most consistent performance across tasks.

6. Conclusion & Discussion

In this work, we introduced **PaCX-MAE**, a framework for distilling latent physiological signals captured via ECG and laboratory values into a frozen visual encoder. By implementing a multimodal contrastive and regression-based objective, we demonstrated that “phantom” physiological priors can be effectively embedded into standard chest X-ray representations without requiring paired data at inference time. Our extensive evaluation across 9 benchmarks reveals that PaCX not only matches state-of-the-art domain-specific baselines on structural tasks like segmentation but significantly outperforms them on physiology-dense diagnostic tasks. Notably, these gains are most prominent in low-data regimes, suggesting that physiological grounding acts as a powerful regularizer, guiding the model to attend to clinically relevant soft-tissue structures rather than spurious correlations.

Limitations include our reliance on single-center data (MIMIC-IV), which restricts phenotypic diversity; multi-center validation is essential to ensure these priors generalize across populations. Additionally, our global alignment strategy overlooks dense, region-specific correlations, such as mapping specific ECG waveforms to localized cardiac sub-regions. Future work will explore granular token-level distillation and longitudinal modeling to capture disease progression and anatomical nuance more holistically.

7. Code Availability

The code for PaCX-MAE is available on GitHub at <https://github.com/Lyce24/PACX-MAE/>. This repository includes the complete implementation and scripts needed to reproduce our results.

References

- Adel, M. Chestx6: Multi-class x-ray dataset. Kaggle, 2025.
- Ahishali, M., Degerli, A., Yamac, M., Kiranyaz, S., Chowdhury, M. E. H., Hameed, K., Hamid, T., Mazhar, R., and Gabbouj, M. Advance warning methodologies for covid-19 using chest x-ray images. *IEEE Access*, 9:41052–41065, 2021. doi: 10.1109/ACCESS.2021.3064927.
- Azizi, S., Mustafa, B., Ryan, F., Beaver, Z., Freyberg, J., Deaton, J., Loh, A., Karthikesalingam, A., Kornblith, S., Chen, T., Natarajan, V., and Norouzi, M. Big self-supervised models advance medical image classification, 2021. URL <https://arxiv.org/abs/2101.05224>.
- Bhave, S., Rodriguez, V., Poterucha, T., Mutasa, S., Aberle, D., Capaccione, K. M., Chen, Y., Dsouza, B., Dumeer, S., Goldstein, J., Hodes, A., Leb, J., Lungren, M., Miller, M., Monoky, D., Navot, B., Wattamwar, K., Wattamwar, A., Clerk, K., Ouyang, D., Ashley, E., Topkara, V. K., Maurer, M., Einstein, A. J., Uriel, N., Homma, S., Schwartz, A., Jaramillo, D., Perotte, A. J., and Elias, P. Deep learning to detect left ventricular structural abnormalities in chest x-rays. *European Heart Journal*, 45(22): 2002–2012, 2024.
- Boecking, B., Usuyama, N., Bannur, S., Castro, D. C., Schwaighofer, A., Hyland, S., Wetscherek, M., Naumann, T., Nori, A., Alvarez-Valle, J., Poon, H., and Oktay, O. *Making the Most of Text Semantics to Improve Biomedical Vision-Language Processing*, pp. 1–21. Springer Nature Switzerland, 2022. ISBN 9783031200595. doi: 10.1007/978-3-031-20059-5_1. URL http://dx.doi.org/10.1007/978-3-031-20059-5_1.
- Candemir, S., Jaeger, S., Palaniappan, K., Musco, J. P., Singh, R. K., Xue, Z., Karargyris, A., Antani, S., Thoma, G., and McDonald, C. J. Lung segmentation in chest radiographs using anatomical atlases with nonrigid registration. *IEEE Transactions on Medical Imaging*, 33(2): 577–590, 2014. doi: 10.1109/TMI.2013.2290491.
- Cho, K., Kim, K. D., Nam, Y., Jeong, J., Kim, J., Choi, C., Lee, S., Lee, J. S., Woo, S., Hong, G.-S., Seo, J. B., and Kim, N. Chess: Chest x-ray pre-trained model via self-supervised contrastive learning. *Journal of Digital Imaging*, 36:902–910, 2023.
- Chowdhury, M. E. H., Rahman, T., Khandakar, A., Mazhar, R., Kadir, M. A., Mahbub, Z. B., Islam, K. R., Khan, M. S., Iqbal, A., Emadi, N. A., Reaz, M. B. I., and Islam, M. T. Can ai help in screening viral and covid-19 pneumonia? *IEEE Access*, 8:132665–132676, 2020. doi: 10.1109/ACCESS.2020.3010287.
- Degerli, A., Ahishali, M., Kiranyaz, S., Chowdhury, M. E. H., and Gabbouj, M. Reliable covid-19 detection using chest x-ray images. In *2021 IEEE International Conference on Image Processing (ICIP)*, pp. 185–189, 2021a. doi: 10.1109/ICIP42928.2021.9506442.
- Degerli, A., Ahishali, M., Yamac, M., Kiranyaz, S., Chowdhury, M. E. H., Hameed, K., Hamid, T., Mazhar, R., and Gabbouj, M. Covid-19 infection map generation and detection from chest x-ray images. *Health Inf Sci Syst.*, 15, 2021b.
- Degerli, A., Kiranyaz, S., Chowdhury, M. E. H., and Gabbouj, M. Osegnet: Operational segmentation network for covid-19 detection using chest x-ray images. In *2022 IEEE International Conference on Image Processing (ICIP)*, pp. 2306–2310, 2022. doi: 10.1109/ICIP46576.2022.9897412.
- Dou, Q., Liu, Q., Heng, P.-A., and Glocker, B. Unpaired multi-modal segmentation via knowledge distillation. *IEEE TMI*, 2020.
- Elias, P. and Bhave, S. Chexchonet: A chest radiograph dataset with gold standard echocardiography labels. PhysioNet, 2024.
- Elsharief, S., Shurrob, S., Jorf, B. A., Lopez, L. J. L., Geras, K. J., and Shamout, F. E. Medmod: Multimodal benchmark for medical prediction tasks with electronic health records and chest x-ray scans. In Xu, X. O., Choi, E., Singhal, P., Gerych, W., Tang, S., Agrawal, M., Subbaswamy, A., Sizikova, E., Dunn, J., Daneshjou, R., Sarker, T., McDermott, M., and Chen, I. (eds.), *Proceedings of the sixth Conference on Health, Inference, and Learning*, volume 287 of *Proceedings of Machine Learning Research*, pp. 781–803. PMLR, 25–27 Jun 2025. URL <https://proceedings.mlr.press/v287/elsharief25a.html>.
- Gorade, V., Sing, A., and Mishra, D. Otcxr: Rethinking self-supervised alignment using optimal transport for chest x-ray analysis. In *2025 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pp. 7143–7152, 2025. doi: 10.1109/WACV61041.2025.00694.
- Gupta, A., Osman, I., Shehata, M. S., and Braun, J. W. Medmae: A self-supervised backbone for medical imaging tasks, 2024. URL <https://arxiv.org/abs/2407.14784>.
- Gupta, S., Hoffman, J., and Malik, J. Cross modal distillation for supervision transfer, 2015. URL <https://arxiv.org/abs/1507.00448>.
- Huang, S.-C., Pareek, A., Jensen, M., Lungren, M. P., Yenung, S., and Chaudhari, A. S. Self-supervised learning

- for medical image classification: a systematic review and implementation guidelines. *npj Digital Medicine*, 6(1):74, Apr 2023. ISSN 2398-6352. doi: 10.1038/s41746-023-00811-0. URL <https://doi.org/10.1038/s41746-023-00811-0>.
- Indeewara, W., Hennayake, M., Rathnayake, K., Ambegoda, T., and Meedeniya, D. Chest x-ray dataset with lung segmentation. PhysioNet, 2023.
- Irvin, J., Rajpurkar, P., Ko, M., Yu, Y., Ciurea-Ilcus, S., Chute, C., Marklund, H., Haghgoo, B., Ball, R., Shpanskaya, K., Seekins, J., Mong, D. A., Halabi, S. S., Sandberg, J. K., Jones, R., Larson, D. B., Langlotz, C. P., Patel, B. N., Lungren, M. P., and Ng, A. Y. CheXpert: A large chest radiograph dataset with uncertainty labels and expert comparison, 2019. URL <https://arxiv.org/abs/1901.07031>.
- Jaeger, S., Karargyris, A., Candemir, S., Siegelman, J., Folio, L., Antani, S., Thoma, G., and McDonald, C. J. Automatic screening for tuberculosis in chest radiographs: a survey. *Quantitative Imaging in Medicine and Surgery*, 3(2), 2013. ISSN 2223-4306. URL <https://qims.amegroups.org/article/view/1813>.
- Jaeger, S., Candemir, S., Antani, S., Wang, Y.-X. J., Lu, P.-X., and Thoma, G. Two public chest x-ray datasets for computer-aided screening of pulmonary diseases. *Quant. Imaging Med. Surg.*, 4(6):475–477, December 2014.
- Li, J., Aguirre, A. D., Junior, V. M., Jin, J., Liu, C., Zhong, L., Sun, C., Clifford, G., Westover, M. B., and Hong, S. An electrocardiogram foundation model built on over 10 million recordings. *NEJM AI*, 2(7):A1oa2401033, 2025. doi: 10.1056/A1oa2401033. URL <https://ai.nejm.org/doi/full/10.1056/A1oa2401033>.
- Lopez-Paz, D., Bottou, L., Schölkopf, B., and Vapnik, V. Unifying distillation and privileged information, 2016. URL <https://arxiv.org/abs/1511.03643>.
- Nguyen, H. Q., Pham, H. H., le tuan linh, Dao, M., and lam khanh. Chest x-ray dataset with lung segmentation. PhysioNet, 2021.
- Nguyen, H. Q., Lam, K., Le, L. T., Pham, H. H., Tran, D. Q., Nguyen, D. B., Le, D. D., Pham, C. M., Tong, H. T. T., Dinh, D. H., Do, C. D., Doan, L. T., Nguyen, C. N., Nguyen, B. T., Nguyen, Q. V., Hoang, A. D., Phan, H. N., Nguyen, A. T., Ho, P. H., Ngo, D. T., Nguyen, N. T., Nguyen, N. T., Dao, M., and Vu, V. Vindr-cxr: An open dataset of chest x-rays with radiologist’s annotations, 2022. URL <https://arxiv.org/abs/2012.15029>.
- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., and Sutskever, I. Learning transferable visual models from natural language supervision, 2021. URL <https://arxiv.org/abs/2103.00020>.
- Rahman, T., Khandakar, A., Qiblawey, Y., Tahir, A., Kiranyaz, S., Kashem, S. B. A., Islam, M. T., Maadeed, S. A., Zughraier, S. M., Khan, M. S., and Chowdhury, M. E. Exploring the effect of image enhancement techniques on covid-19 detection using chest x-ray images. *Computers in Biology and Medicine*, 132:104319, 2021.
- Saporta, A., Puli, A., Goldstein, M., and Ranganath, R. Contrasting with symile: Simple model-agnostic representation learning for unlimited modalities, 2024. URL <https://arxiv.org/abs/2411.01053>.
- Saporta, A., Puli, A. M., Goldstein, M., and Ranganath, R. Symile-mimic: a multimodal clinical dataset of chest x-rays, electrocardiograms, and blood labs from mimic-iv. PhysioNet, 2025.
- Tahir, A. M., Chowdhury, M. E. H., Qiblawey, Y., Khandakar, A., Rahman, T., Kiranyaz, S., Khurshid, U., Ibtehaz, N., Mahmud, S., and Ezeddin, M. Covid-qu-ex. Kaggle, 2021. URL <https://doi.org/10.34740/kaggle/dsv/3122958>.
- Tiu, E., Talius, E., Patel, P., Langlotz, C. P., Ng, A. Y., and Rajpurkar, P. Expert-level detection of pathologies from unannotated chest x-ray images via self-supervised learning. *Nature Biomedical Engineering*, 6(12):1399–1406, Dec 2022. ISSN 2157-846X. doi: 10.1038/s41551-022-00936-9. URL <https://doi.org/10.1038/s41551-022-00936-9>.
- Wang, H. et al. Learnable cross-modal knowledge distillation for multi-modal learning with missing modality. *arXiv preprint arXiv:2310.01035*, 2023.
- Wang, X., Peng, Y., Lu, L., Lu, Z., Bagheri, M., and Summers, R. M. Chestx-ray8: Hospital-scale chest x-ray database and benchmarks on weakly-supervised classification and localization of common thorax diseases. *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 3462–3471, 2017. URL <https://api.semanticscholar.org/CorpusID:263796294>.
- Xiao, J., Bai, Y., Yuille, A., and Zhou, Z. Delving into masked autoencoders for multi-label thorax disease classification. In *2023 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pp. 3577–3589, 2023. doi: 10.1109/WACV56688.2023.00358.

Yamaç, M., Ahishali, M., Degerli, A., Kiranyaz, S., Chowdhury, M. E. H., and Gabouj, M. Convolutional sparse support estimator-based covid-19 recognition from x-ray images. *IEEE Transactions on Neural Networks and Learning Systems*, 32(5):1810–1820, 2021. doi: 10.1109/TNNLS.2021.3070467.

Zhang, Y., Jiang, H., Miura, Y., Manning, C. D., and Langlotz, C. P. Contrastive learning of medical visual representations from paired images and text, 2022. URL <https://arxiv.org/abs/2010.00747>.

Zhou, L., Liu, H., Bae, J., He, J., Samaras, D., and Prasanna, P. Self pre-training with masked autoencoders for medical image classification and segmentation, 2023. URL <https://arxiv.org/abs/2203.05573>.

A. Appendix

A.1. Methodology

A.1.1. UNIMODAL ENCODER PRETRAINING

CheXpert Dataset (Pretraining Data). We initialize the visual backbone using the CheXpert dataset (Irvin et al., 2019), a large-scale collection of chest radiographs. The training split consists of approximately **224k** frontal and lateral chest X-rays. All images are resized to 224×224 pixels and normalized using standard ImageNet statistics (mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]) prior to patch embedding. No patient labels are utilized during this stage; the objective is purely self-supervised reconstruction.

CXR Encoder (Masked Autoencoder). We employ a Vision Transformer (ViT-B/16) architecture. The input image $x \in \mathbb{R}^{224 \times 224 \times 3}$ is divided into non-overlapping patches of size 16×16 , resulting in a sequence of $N = 196$ tokens.

- **Masking Strategy:** We employ an aggressive random masking ratio of **0.90** (90% of patches are discarded), significantly higher than the standard 0.75 used in natural image MAE (Gupta et al., 2024), to prevent the model from relying on local interpolation of smooth tissues.
- **Architecture:** The encoder operates only on the visible set of patch embeddings (latent dim $D = 768$, depth=12, heads=12). A lightweight decoder (depth=8, dim=512, heads=16) reconstructs the pixel values of the masked patches.
- **Objective:** We minimize the Mean Squared Error (MSE) between the reconstructed and original patches:

$$\mathcal{L}_{\text{MAE}} = \frac{1}{|\mathcal{M}|} \sum_{i \in \mathcal{M}} \|\hat{x}_i - x_i\|_2^2 \quad (1)$$

where \mathcal{M} is the set of masked indices.

- **Optimization:** The model is trained for 400 epochs using the **AdamW** optimizer ($\beta_1 = 0.9$, $\beta_2 = 0.95$, weight decay 0.05). The learning rate follows a schedule with a base value of 1.5×10^{-4} , featuring a linear warmup for 10 epochs followed by a cosine decay to zero. The weight decay is 0.05.

Laboratory Encoder (Denoising Autoencoder). We train a modality-specific encoder on the **Symile-MIMIC dataset** to handle the sparse, tabular nature of laboratory measurements.

- **Data Representation:** The input consists of N laboratory tests converted to percentiles $x \in [0, 1]^N$ and a binary missingness mask $m \in \{0, 1\}^N$. In our case, $N = 50$.
- **Architecture:** The model is a symmetric MLP. The input layer concatenates values and masks ($x \oplus m \in \mathbb{R}^{2N}$). The encoder projects this to a hidden layer of size 512 (with LayerNorm and GELU activation) and then to a bottleneck latent representation $z \in \mathbb{R}^{256}$. The decoder mirrors this structure ($256 \rightarrow 512 \rightarrow N$), concluding with a Sigmoid activation to ensure outputs remain in the valid percentile range $[0, 1]$.
- **Denoising Objective:** We employ a dynamic corruption strategy where **15%** of the *observed* values in x are randomly zeroed out during training. The model is trained to reconstruct the original values minimizing the MSE only on the valid indices:

$$\mathcal{L}_{\text{DAE}} = \frac{\sum_{j=1}^N m_j \cdot (\hat{x}_j - x_j)^2}{\sum_{j=1}^N m_j + \epsilon} \quad (2)$$

- **Optimization:** We train for 50 epochs (batch size 256) using **AdamW** (learning rate 3×10^{-4} , weight decay 0.01). We use a step-based scheduler with a linear warmup of **200 steps** (approx. 1 epoch) followed by cosine decay.

A.1.2. CROSS-MODAL DISTILLATION LEARNING

Parameter-Efficient Fine-Tuning (PEFT). To align the CXR encoder (f_v) with the frozen physiological targets (f_e, f_l), we freeze the entire ViT backbone except for:

1. **Normalization Layers:** All ‘LayerNorm’ parameters are unfrozen to adapt to the domain shift from CheXpert to MIMIC-CXR.
2. **LoRA Modules:** We inject Low-Rank Adaptation matrices ($r = 8, \alpha = 8$) into the Query, Key, Value (‘qkv’) and Feed-Forward (‘fc’) projection layers of every transformer block. This accounts for < 1% of total trainable parameters.

Dual Objective Details. We train with a weighted sum of contrastive and regression losses: $\mathcal{L}_{total} = 0.5 \cdot \mathcal{L}_{CLIP} + 1.0 \cdot \mathcal{L}_{Reg}$. We choose the loss weights based on the loss scales.

(i) Cross-modal CLIP Alignment. We map the CXR (h_{cxr}) and physiological (z_m) embeddings to a shared 256-dimensional space via linear projection heads. We compute the symmetric InfoNCE loss:

$$\ell_{clip}(u, v) = -\log \frac{\exp(u^\top v / \tau)}{\sum_{k=1}^B \exp(u^\top v_k / \tau)} \quad (3)$$

Key implementation details:

- **Temperature (τ):** initialized at 0.07 and learned during training.
- **Global Negatives:** We utilize ‘all_gather’ to collect embeddings from all available GPUs, effectively scaling the number of negative samples by the number of devices (e.g., $B_{eff} = B \times N_{GPUs}$).
- **Label Smoothing:** We apply $\epsilon = 0.02$ smoothing to the target distribution to mitigate overfitting to noisy clinical pairings.

(ii) Physiology Prediction (Regression). We employ modality-specific regression heads (Linear $768 \rightarrow D_{target}$) to predict the raw, unprojected output of the frozen physiological encoders. The loss is the Cosine Distance:

$$\mathcal{L}_{reg} = \sum_{m \in \{ecg, lab\}} \frac{1}{2} \left(1 - \frac{r_m(h_{cxr}) \cdot \text{sg}[z_m]}{\|r_m(h_{cxr})\| \|\text{sg}[z_m]\|} \right) \quad (4)$$

where $\text{sg}[\cdot]$ indicates the stop-gradient operator, ensuring the physiological encoders remain purely frozen targets.

A.2. Experiments

A.2.1. DATASETS

We evaluate PaCX-MAE on a diverse set of public chest X-ray benchmarks spanning classification and segmentation tasks.

Classification Benchmarks

- **TB:** A dataset provided by the National Library of Medicine (NIH) containing 662 CXRs labeled as tuberculosis or non-tuberculosis (Jaeger et al., 2014). The dataset is balanced (326 normal, 336 TB). As no official split is provided, we use a 70%/10%/20% train/validation/test split.
- **VinDr-CXR:** A large-scale benchmark of 18,000 CXRs annotated with up to 28 disease labels (Nguyen et al., 2021). We exclude labels with fewer than 50 occurrences and follow the official train/test split, holding out 10% of the training set for validation.
- **NIH-14:** A standard benchmark consisting of 112,120 CXRs from 30,805 patients, annotated with 14 thoracic disease labels (Wang et al., 2017). We follow the official split, using 15% of the training data for validation.
- **ChestX6:** A dataset of 18,036 CXRs labeled with six chest conditions (Adel, 2025). We use the official training, validation, and testing splits.
- **CheXchoNet:** Contains 71,589 CXRs from 24,689 patients, annotated with a composite label indicating severe left ventricular hypertrophy or dilated left ventricle (Elias & Bhave, 2024). As no official split is provided, we adopt a 70%/10%/20% split.
- **MedMod:** A multimodal benchmark derived from MIMIC-IV and MIMIC-CXR (Elsharief et al., 2025). We use a subset of 9,098 CXRs from 8,035 patients (excluding those overlapping with our Symile-MIMIC pretraining data), annotated with 28 binary diagnostic categories spanning cardiopulmonary, circulatory, gastrointestinal, and endocrine conditions. We apply a 70%/10%/20% split.

Segmentation Benchmarks

- **CXLSeg**: A large-scale lung segmentation dataset comprising 243,324 CXRs from MIMIC-CXR with corresponding lung masks (Indeewara et al., 2023). We follow the official training, validation, and testing splits.
- **COVID-QU-Ex (Segmentation)**: A dataset of 33,920 CXRs with segmentation masks for lung and infection regions (Tahir et al., 2021). Images are categorized as COVID-positive, non-COVID infection, or COVID-negative. We use the official splits.
- **QaTa-COV19**: A COVID-19 pneumonia segmentation dataset containing 9,258 CXRs with infection masks. As no official split is provided, we use a 70%/10%/20% train/validation/test split.

A.2.2. IMPLEMENTATION DETAILS

Downstream Classification For standard downstream classification, we employ a linear probing protocol: a linear classifier is trained on top of the frozen backbone for 40 epochs using the Adam optimizer with a learning rate of $\eta = 3 \times 10^{-3}$ and a cosine decay schedule. Performance is reported via AUROC and F1-score. We used effective batch sizes of 512 distributed across GPUs. To simulate low-data regimes, we trained on stratified subsets of the training data (e.g., 10%), monitoring validation AUROC (macro or binary) to checkpoint the best-performing models.

Semantic Segmentation For segmentation, we attach a lightweight 5-layer decoder to the encoder backbone. Models are trained for 50 epochs ($\eta = 1 \times 10^{-4}$) using a composite objective summing Dice loss and Binary Cross-Entropy: $\mathcal{L}_{\text{seg}} = \mathcal{L}_{\text{Dice}} + \mathcal{L}_{\text{BCE}}$.

Physiological Alignment Evaluation We assess the alignment between image representations and physiological states using the Symile-MIMIC dataset (paired CXR, ECG, and Labs) (Saporta et al., 2025).

- **Regression Probe**: We extract frozen embeddings from the ViT-Base backbone (16 × 16 patch size, 224 × 224 resolution) and train a Ridge regression probe ($\alpha = 10.0$, solver=“svd”) to predict normalized ECG and Lab feature vectors. Performance is measured via R^2 and Cosine Similarity (raw and centered).
- **Cross-Modal Retrieval**: We compute Recall@ K ($K \in \{1, 5, 10\}$) to evaluate the model’s ability to retrieve the correct physiological profile for a given patient from the test set based solely on their chest X-ray.

Ablation Studies We investigate component contributions through two specific protocols:

- **Modality Ablation**: To measure the impact of specific physiological signals, we employ a masking protocol where specific modalities (ECG or Labs) are set to `None` during training. This removes them from both the cross-modal projection mechanism and regression targets.
- **Loss Ablation**: We isolate the contributions of the contrastive and regression objectives by manipulating their scalar weights ($\lambda_{\text{clip}}, \lambda_{\text{reg}}$). By setting a specific coefficient to 0, we neutralize that objective’s contribution to gradient updates while maintaining consistent training dynamics.

Interpretability (Attention Rollout) To visualize model focus, we implement Attention Rollout. We patch the ViT blocks to cache raw attention scores and recursively multiply attention matrices from the input to the final layer, adding residual connections and averaging across heads. The resulting attention map for the [CLS] token is interpolated from the 14 × 14 grid to the original 224 × 224 resolution and min-max normalized for visualization.

A.3. Results

A.3.1. ADDITIONAL CASES FOR ATTENTION ROLLOUT

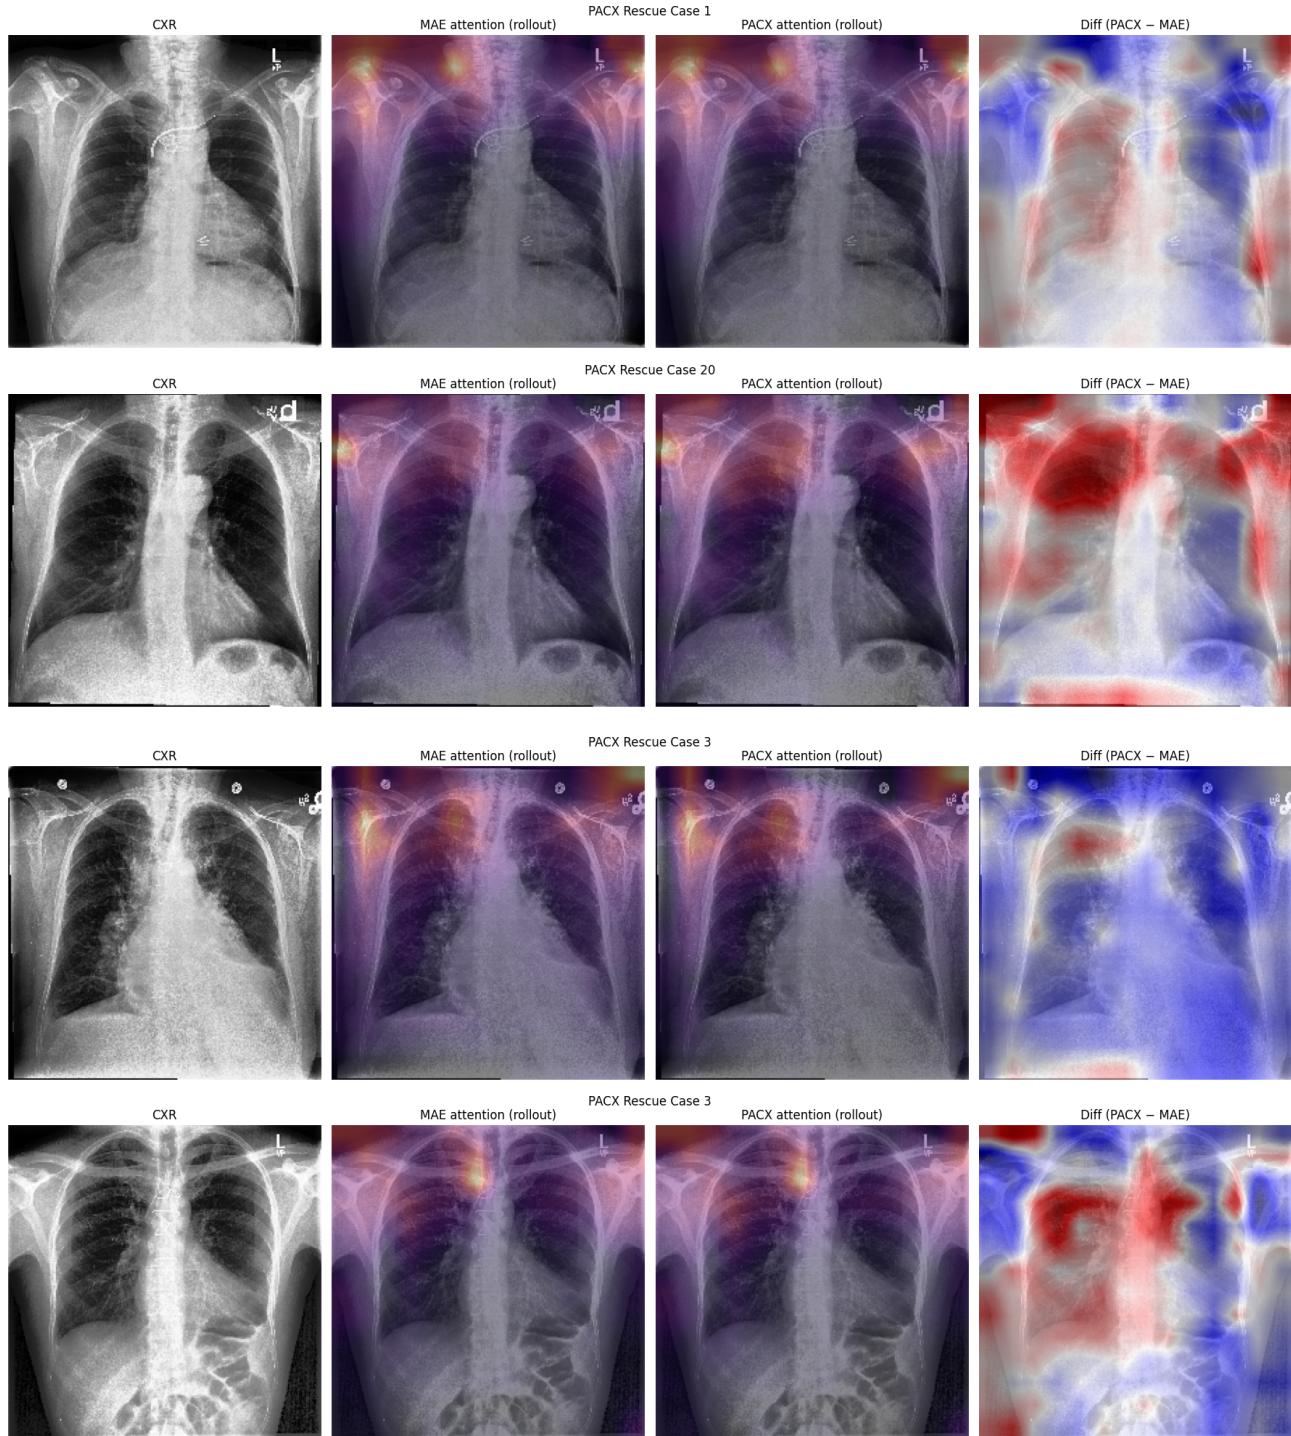


Figure 5. Additional Attention Rollout Cases. These visualizations confirm the consistent trend of PaCX attending to soft-tissue anatomy versus the edge-focused attention of the MAE baseline.

B. Acknowledgment

We received assistance from ChatGPT and Gemini throughout this project. This included support with literature studies, data processing, model implementation, and refinement of our manuscript, including L^AT_EX formatting and overall readability.

C. Contribution

Yancheng Liu spearheaded the project’s conceptualization and technical execution. He formulated the core research objective of enabling “multimodal training with unimodal inference” and architected the two-stage PaCX framework. His specific contributions include:

- **Framework Design & Pretraining:** He designed and trained the MAE visual backbone, conducting extensive hyperparameter sweeps on the CheXpert dataset to optimize reconstruction quality. He further engineered the physiological encoding stream by training a Denoising Autoencoder for laboratory values and integrating the ECGFounder architecture.
- **Distillation Strategy:** He led the rigorous architectural search for the cross-modal distillation mechanism. This involved designing and evaluating multiple training strategies (fixed temperature scaling, embedding mixup, teacher-student setups, and partial ViT unfreezing versus LoRA adaptation) to identify the optimal hyperparameters for physiological alignment.
- **Engineering & Implementation:** He developed the unified Pytorch Lightning codebase with Stable—Pretraining support from Rbalestr Lab (<https://github.com/rbalestr-lab/stable-pretraining>) for the MAE, cross-modal distillation, and evaluation stages. He also managed the complete data lifecycle, implementing preprocessing pipelines, conducting statistical sanity checks, and building robust DataModules for all 9 benchmark datasets.
- **Evaluation & Reporting:** He did the comprehensive experimental roadmap, logging approximately 500 GPU hours to execute linear probing, low-data regime analysis, zero-shot alignment testing, and Attention Rollout visualizations. Finally, he authored the Methodology, Results, and Appendix sections of this report.

Kenichi Maeda initially explored the use of the MC-MED dataset, which was eventually abandoned. He was also involved in literature studies, planning of the multimodal injection step (e.g., adaptation of the Symile-mimic dataset and the original Symile model, initial setup of teacher-student style distillation learning, etc.), segmentation evaluation, UMAP visualization, as well as modality and loss ablation studies.

Manan Pancholy initially proposed the use of the MC-MED dataset, which was eventually abandoned. Afterward, he proposed the use of the Symile-MIMIC dataset and initially implemented a SimCLR-adjacent framework that employed cross-modal attention for the multimodal injection, the latter of which was eventually abandoned. He also conducted literature review (real-world clinical applications, medically-inspired data augmentation pipelines, and available datasets), classification evaluation (including preprocessing the CheXchoNet and MedMod datasets), modality ablation studies, and loss ablation studies. He wrote and refined their corresponding portions of the manuscript as well as the discussion section.